


Intelligent Multi Agent Systems



Computing Solution Concepts
Gerhard Neumann

Agenda



Now we know a lot about different solution concepts...

➡ But how to compute them?

➡ We will get to know a few algorithms for different solution concepts

Outline



Algorithms for:

- ➡ Zero Sum Games
- ➡ Computing Nash Equilibria
- ➡ Iterated Dominance

2 Player Zero-Sum Games



Recap: $u_1(a_1, a_2) = -u_2(a_1, a_2)$

➡ Optimal strategy: $U_1^*(\alpha) = -U_2^*(\alpha)$

➡ All Nash Equilibria have the same value $U_1^*(\alpha)$

➡ All Nash Equilibria are minmax and maxmin strategies

$$\arg \max_{\alpha_1} \min_{\alpha_2} u_1(\alpha_1, \alpha_2)$$

$$\arg \min_{\alpha_1} \max_{\alpha_2} u_2(\alpha_1, \alpha_2)$$



2 Player Zero-Sum Games



We can construct a **linear program**, that implements the **minmax strategy** for player 2:

$$\begin{aligned} & \arg \min_{U_1^*, \alpha_2} U_1^* \\ & \text{s.t.} \quad \sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k \leq U_1^*, \quad \forall j \in A_1 \\ & \quad \sum_{k \in A_2} \alpha_2^k = 1 \\ & \quad \alpha_2^k \geq 0, \quad \forall k \in A_2 \end{aligned}$$



MinMax Strategies

$$\arg \min_{U_1^*, \alpha_2} U_1^*$$

$$\text{s.t. } \sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k \leq U_1^*, \quad \forall j \in A_1$$

$$\sum_{k \in A_2} \alpha_2^k = 1$$

$$\alpha_2^k \geq 0, \quad \forall k \in A_2$$

Variables: U_1^*, α_2^k

- ➡ First constraint says: U_1^* is the maximum value for player 1 (in combination with objective) if current strategy of player 2 is α_2^k
- ➡ We want to minimize the maximum value U_1^*
- ➡ That's a **MinMax Strategy** for player 2!



... and vice versa: MaxMin Strategies

$$\arg \max_{U_1^*, \alpha_1} U_1^*$$

$$\text{s.t. } \sum_{j \in A_1} u_1(a_1^j, a_2^k) \alpha_1^j \geq U_1^*, \quad \forall k \in A_2$$

$$\sum_{j \in A_1} \alpha_1^j = 1$$

$$\alpha_1^j \geq 0, \quad \forall j \in A_1$$

Variables: U_1^*, α_1^j

➡ First constraint says: U_1^* is the **worst case value** for player 1 with current strategy α_1^j

➡ We want to maximize the worst-case value U_1^*

➡ That's a **MaxMin Strategy** for player 1!



Useful alternative formulation (MinMax)

Formulation with Slack Variables:

$$\begin{aligned} \min \quad & U_1^* \\ \text{s.t.} \quad & \sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k + r_1^j = U_1^*, \quad \forall j \in A_1 \\ & \sum_{k \in A_2} \alpha_2^k = 1 \\ & \alpha_2^k \geq 0, \quad \forall k \in A_2, \quad r_1^j \geq 0, \quad \forall j \in A_1 \end{aligned}$$

Variables: U_1^*, α_2^k, r_1^j

➡ Slack variables must all be positive

➡ Therefore, the equality constraint is equal to the inequality constraint from the previous slides

Outline



Algorithms for:

- ➡ Zero Sum Games
- ➡ Computing Nash Equilibria
- ➡ Iterated Dominance

2 Player General Sum Games



Linear Program's are nice: Solving LP's is in P

Unfortunately, we can not formulate a linear program any more

- ➡ No opposing interests
- ➡ Can not minimize utility of other agent to maximize own utility
- ➡ Yet, General Sum Games can be formulated as **Linear Complementarity Problem (LCP)**
- ➡ Solving an LCP is PPAD complete: “Polynomial parity argument directed version” (almost as hard as NP-complete)

2 Player General Sum Games



General Sum Games can be formulated as **Linear Complementarity Problem (LCP)**

$$\sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k + r_1^j = U_1^*, \quad \forall j \in A_1$$

$$\sum_{j \in A_1} u_2(a_1^j, a_2^k) \alpha_1^j + r_2^k = U_2^*, \quad \forall k \in A_2$$

$$\sum_{j \in A_1} \alpha_1^j = 1, \quad \sum_{k \in A_2} \alpha_2^k = 1$$

$$\alpha_1^j \geq 0, \quad \forall j \in A_1, \quad r_2^k \geq 0, \quad \forall k \in A_2$$

$$\alpha_2^k \geq 0, \quad \forall k \in A_2, \quad r_1^j \geq 0, \quad \forall j \in A_1$$

$$r_1^j \alpha_1^j = 0, \quad r_2^k \alpha_2^k = 0, \quad \forall j \in A_1, \quad \forall k \in A_2$$

2 Player General Sum Games



General Sum Games can be formulated as **Linear Complementarity Problem (LCP)**

$$\sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k + r_1^j = U_1^*, \quad \forall j \in A_1$$

$$\sum_{j \in A_1} u_2(a_1^j, a_2^k) \alpha_1^j + r_2^k = U_2^*, \quad \forall k \in A_2$$

$$\sum_{j \in A_1} \alpha_1^j = 1, \quad \sum_{k \in A_2} \alpha_2^k = 1$$

$$\alpha_1^j \geq 0, \quad \forall j \in A_1, \quad r_2^k \geq 0, \quad \forall k \in A_2$$

$$\alpha_2^k \geq 0, \quad \forall k \in A_2, \quad r_1^j \geq 0, \quad \forall j \in A_1$$

$$r_1^j \alpha_1^j = 0, \quad r_2^k \alpha_2^k = 0, \quad \forall j \in A_1, \quad \forall k \in A_2$$

Differences to Zero-Sum Case:

- ➔ No objective, only constraints
- ➔ Constraints for both players in the formulation

Complementary condition:

- ➔ Slack variables always positive
- ➔ If action a_1^j is used (i.e. $\alpha_1^j > 0$) then $r_1^j = 0$

Linear Complementarity Problem



$$\sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k + r_1^j = U_1^*, \quad \forall j \in A_1, \quad \sum_{j \in A_1} u_2(a_1^j, a_2^k) \alpha_1^j + r_2^k = U_2^*, \quad \forall k \in A_2$$
$$r_1^j \alpha_1^j = 0, \quad r_2^k \alpha_2^k = 0, \quad \forall j \in A_1, \quad \forall k \in A_2$$

Why does it solve our problem?

- ➡ If action a_1^j is used (i.e. $\alpha_1^j > 0$) then $r_1^j = 0$
- ➡ For each action \bar{a}_1^j with $\alpha_1^j > 0$ it holds:
 - ➡ The expected utility for agent 1 for all actions \bar{a}_1^j is the same
$$\sum_{k \in A_2} u_1(a_1^j, a_2^k) \alpha_2^k = U_1^*, \quad \forall j, \alpha_1^j > 0$$
 - ➡ The value U_1^* is the maximum expected utility agent 1 can get (assuming strategy from agent 2 is α_2)
 - ➡ Each action \bar{a}_1^j is a best response to α_2

Linear Complementarity Problem



The same argument holds for player 2

➡ If the LCP is satisfied, we have found a (mixed strategy) Nash Equilibrium!

Ok, how do we solve a LCP?:

- ➡ There are many algorithms which do that...
- ➡ Most famous one: Lemke Howson Algorithm
- ➡ Only finds one Nash Equilibrium

Computationally very expensive:

- ➡ Solving an LCP is PPAD complete: “Polynomial parity argument directed version”
- ➡ People believe that PPAD is much harder than P (similar to NP)
- ➡ Worst case: Exponential in size of the game (?)
- ➡ But there is no proof... (also similar to NP)

Beyond finding an Equilibrium



We might want to find a equilibrium with one of the following properties:

- ➔ **Uniqueness:** Is there a unique Equilibrium for game G ?
- ➔ **Pareto optimality:** Does there exist a strict Pareto Optimal equilibrium?
- ➔ **Guaranteed Payoff:** Does there exist a equilibrium where some player i gets an expected payoff of at least v ?
- ➔ **Guaranteed social welfare:** Does there exist an equilibrium where the sum of all agent's utility is at least k ?
- ➔ **Action inclusion:** Does there exist a equilibrium where player i plays action a_i with positive probability
- ➔ **Action exclusion:** Does there exist a equilibrium where player i plays action a_i with zero probability

Bad news: All these questions are NP-complete!

Outline



Algorithms for:

- ➡ Zero Sum Games
- ➡ Computing Nash Equilibria
- ➡ Iterated Dominance

Identifying dominated strategies



We can make the problem easier by **deleting dominated strategies**:

Definition:

In a strategic game player i 's action a_i'' **strictly dominates** another action a_i' if

- $u_i(a_i'', a_{-i}) > u_i(a_i', a_{-i})$ for every list a_{-i} of the other player's action

We say that a_i' is **strictly dominated**

Definition:

In a strategic game player i 's action a_i'' **weakly dominates** another action a_i' if

- $u_i(a_i'', a_{-i}) \geq u_i(a_i', a_{-i})$ for **every** list a_{-i} of the other player's action
- $u_i(a_i'', a_{-i}) > u_i(a_i', a_{-i})$ for **some** list a_{-i} of the other player's action

We say that a_i' is **weakly dominated**

Identifying dominated strategies



However, actions can also be **dominated by mixed strategies**:

- ➡ M is not dominated by U or D
- ➡ But M is dominated by **a mixed strategy** that takes D and U with equal probability.

	L	C
U	3, 1	0, 1
M	1, 1	1, 1
D	0, 1	4, 1

Detecting Dominated Strategies



Test for detecting **strictly dominated** strategies a_i^j :

$$\sum_{k \in A_i} \alpha_i^k u_i(a_i^k, a_{-i}) > u_i(a_i^j, a_{-i}) \quad \forall a_{-i} \in A_{-i}$$

$$\alpha_i^k \geq 0, \quad \forall k \in A_i$$

$$\sum_{k \in A_i} \alpha_i^k = 1$$

There exists an α_i such that the expected utility of α_i is always larger than $u_i(a_i^j, a_{-i})$, no matter what the other agents do.

- ➡ However, this not a proper linear program
- ➡ No objective, we need weak inequalities
- ➡ We will also assume that all utilities are positive (> 0)

Detecting dominated strategies



Linear Program for detecting **strictly dominated strategies**:

$$\begin{aligned} & \text{minimize} && \sum_{k \in A_i} \alpha_i^k \\ & \text{subject to} && \sum_{k \in A_i} \alpha_i^k u_i(a_i^k, a_{-i}) \geq u_i(a_i^j, a_{-i}) && \forall a_{-i} \in A_{-i} \\ & && \alpha_i^k \geq 0 && \forall k \in A_i \end{aligned}$$

- ➡ No normalizing constraint
- ➡ Minimization of summed “probabilities”

Why is it a solution to our problem?

- ➡ If a solution exists with $\sum_k \alpha_i^k < 1$ then we can add $1 - \sum_k \alpha_i^k$ to some α_i^k and we'll have a dominating mixed strategy (since utility was assumed to be positive everywhere)

Iterated Dominance



Similar programs can be constructed for **weakly dominated strategies**.

This can be done by **repeatedly** solving our LPs:

- ➡ By deleting strategies, other strategies might become dominated.
- ➡ Checking whether every pure strategy of every player is dominated by any other mixed strategy requires us to solve at worst $\sum_{i \in N} |A_i|$ linear programs.
- ➡ Each step removes one pure strategy for one player, so there can be at most $\sum_{i \in N} (|A_i| - 1)$ steps.
- ➡ Thus we need to solve $O((n \cdot \max_i |A_i|)^2)$ linear programs.

We might reduce the action sets considerably, simplifying the use of LCP

Outline



Algorithms for:

- ➡ Zero Sum Games
- ➡ Computing Nash Equilibria
- ➡ Iterated Dominance